

BUSINESS MODEL & STRATEGIC BLUEPRINT

RTFCLMGZN

artificial magazine

An AI-native news publication, written by an AI editorial staff, about the AI industry itself.

Prepared: July 8, 2026

Confidential — internal planning document, for founder use in direction and deployment.

Contents

1. Executive Summary
2. The Concept
3. Market Landscape & Competitive Analysis
4. Content Strategy & Topic Allocation
5. The AI Newsroom — Editorial Agent System
6. Product & Platform
7. Monetization Strategy & Pricing Tiers
8. Technology Stack & Operating Cost Structure
9. Podcast Strategy
10. Social Media Strategy
11. Legal, Compliance & Risk
12. Illustrative Financial Model
13. Phased Rollout Plan
14. Appendix — Persona Sheet & Pipeline Glossary

1. Executive Summary

RTFCLMGZN is a proposed news and magazine business built on a specific structural bet: the AI-news category today is dominated by human-run email newsletters — The Rundown AI, TLDR AI, Superhuman AI, Ben's Bites, Import AI, and a long tail of smaller ones. All of them are essentially the same product: a daily or weekly digest email, written by one or two named humans, monetized primarily through sponsor placements. No player in the category has yet built a full website-based publication with a distinctive, multi-personality AI editorial staff, an original-article model rather than a pure-digest model, a companion podcast, and a premium print-style weekly issue — where the news operation itself is produced end-to-end by a coordinated system of AI agents.

That gap is the opportunity. The audience already exists and is large and growing — AI newsletter subscriptions grew 340% between 2023 and 2025, and the broader AI-generated newsletter market is projected to reach roughly \$2.5 billion in 2026 and \$6.7 billion by 2030. But “AI-obsessed” readers are also the audience most receptive to, and most likely to talk about, a publication that is transparently produced by AI itself — that premise is the built-in marketing hook, not a liability to hide.

This blueprint lays out: the concept and positioning, a competitive read of the newsletter landscape, a content and topic strategy (roughly 85% technology-broad coverage anchored in AI, with small deliberate slices of health, crypto, world news, and quantum), a proposed AI “newsroom” architecture with named editor personas and a meta agent group that designs and staffs the operation, a product and monetization plan with concrete pricing tiers, a cost model built on current API pricing for text, image, and voice generation, a legal and risk checklist specific to AI-authored news, an illustrative financial model, and a phased rollout plan.

Treat every number in this document as a planning assumption to pressure-test, not a guarantee. Pricing, subscriber, and revenue figures are illustrative estimates built from current public data (cited throughout) and industry benchmarks, not projections of RTFCLMGZN's actual performance. This is not legal, tax, or financial advice — validate the business, tax, and compliance decisions with a qualified lawyer and accountant before launch.

2. The Concept

2.1 The name

RTFCLMGZN is a vowel-stripped compression of “ARTIFICIAL MAGAZINE.” The joke is also the thesis, stated three times over: it is an artificial (AI-built) magazine, produced by an artificial staff, reporting on artificial intelligence. That triple-layer pun is memorable, unmistakably on-theme for an AI-obsessed audience, and — because it reads like a compressed file name or a variable — it visually signals “built by a machine” before a reader has read a single word.

2.2 The elevator pitch

A daily-updated website covering AI and the technology industry around it, written and produced entirely by a coordinated system of AI editorial agents, each with a distinct beat, voice, and personality. Free to read. Extended by a twice-weekly podcast, a paid weekly print-style digital magazine, and active social distribution — all generated by the same agent system, overseen by a lightweight human review layer for accuracy and liability-sensitive topics.

2.3 What makes this different from the existing field

- **Website-first, not inbox-first.** Every major competitor is an email product. A website supports search, archives, categorization, embeds, a paid tier, and SEO-driven discovery that an email digest cannot.
- **A newsroom, not a newsletter voice.** Competitors run on one named human voice (Rowan Cheung, Ben Tossell, Jack Clark). RTFCLMGZN's pitch is a full masthead of distinct AI personas, each publishing under their own byline, beat, and tone — closer to a real magazine's structure than a solo newsletter.
- **Broader than pure AI.** Competitors cover AI almost exclusively. RTFCLMGZN treats AI as the gravitational center of a wider technology beat, with deliberate, capped slices of health, crypto, world/policy news, and quantum computing — subjects an AI-obsessed reader also tracks, but currently has to get from separate newsletters.
- **Multi-format from day one.** Article, podcast, and paid print-style magazine are built as one content pipeline with three outputs, not three separate businesses bolted together later.
- **The production model is the story.** An entirely AI-run editorial operation, honestly disclosed, is itself a piece of news that an AI-obsessed readership is primed to find interesting and to share — turning the operating model into free distribution.

3. Market Landscape & Competitive Analysis

3.1 The current field

The dominant AI-news publications in mid-2026 are all free, sponsor-funded email newsletters. The table below summarizes the largest ones.

Publication	Subscribers	Format	Monetization
The Rundown AI	~2M+	Free daily email, ~5 min read	Sponsor placements only; separate paid “Rundown University” course at \$84/mo
TLDR AI	~1.1–1.25M	Free daily email, technical/dense	Sponsor placements (part of the wider TLDR network)
Superhuman AI	~1.5–2M	Free, 3-minute daily read	Sponsor placements (AWS, Google, Okta have sponsored)
Ben's Bites	~163–400K	Free daily + optional paid tier	Sponsorships; Pro tier ~\$150/yr for deeper research content
Import AI	100K+	Free weekly, research/policy depth	Independent, written by Jack Clark (Anthropic co-founder)

Figures compiled from multiple 2026 newsletter-landscape reviews (Readless, DataCamp, Dupple, ClickITTech, DemandSage); subscriber counts are publisher-stated and should be treated as approximate.

3.2 Category size and trajectory

AI newsletter subscriptions grew roughly 340% between 2023 and 2025, comfortably outpacing every other newsletter category. The broader AI-generated newsletter market is estimated at about \$2.53 billion in 2026, projected to reach \$6.74 billion by 2030. Reader behavior is also shifting toward consolidation: several 2026 reviews note that the top five AI newsletters overlap on roughly 80% of their lead stories on any major launch day, and third-party “digest of digests” tools have emerged specifically to deduplicate that overlap for readers — a strong signal that the newsletter format has hit a redundancy ceiling and that readers are actively looking for a single, deeper, less repetitive source.

3.3 The whitespace

- **No website-native destination.** Every major competitor optimizes for the inbox. A searchable, categorized, embeddable website with a real archive is structurally absent from the top of the category.
- **No multi-voice masthead.** Every competitor is a single-author (or single-brand) voice. Nobody is running a newsroom model with differentiated beats and personalities.
- **No paid print-style artifact.** Ben's Bites and Rundown both monetize education products (community, courses) rather than a weekly premium issue in the traditional magazine sense.
- **No integrated audio-first extension.** Latent Space pairs a newsletter with a podcast, but as an engineering-audience niche play, not a general AI-news brand.

- **Consolidation fatigue is a marketing opening.** A publication that positions itself explicitly as “one deep, well-sourced read instead of five overlapping ones” speaks directly to the frustration reviewers are already describing in 2026.

4. Content Strategy & Topic Allocation

4.1 Proposed topic mix

Your notes specify technology broadly (not exclusively AI) at roughly 85%, with small, deliberate slices of health, crypto, world news, and quantum layered on top. The table below turns that into a working allocation that sums to 100% and keeps AI itself as the clear gravitational center within the 85% technology figure.

Category	Target share	What it covers
AI industry & models (core)	55%	Model releases, lab news, benchmarks, funding, tools, applied use cases
Broader technology	30%	Robotics, semiconductors/chips, spaceflight, consumer tech, enterprise software — anything adjacent to or shaped by AI
Crypto & digital assets	5%	Where AI and crypto intersect: AI agents transacting on-chain, compute-token markets, major market moves
Health & biotech	4%	AI in medicine, diagnostics, drug discovery, biotech breakthroughs
World news & policy	4%	AI regulation, export controls, geopolitics of compute and chips, major world events with AI relevance
Quantum computing	2%	Hardware milestones, funding, and the quantum-AI intersection

This allocation is a starting editorial budget, not a hard quota — the AI-industry share should flex upward around major news cycles (a frontier model launch, a major acquisition, an export-control change) and the smaller categories exist to keep the publication from reading as a single-topic trade newsletter.

4.2 Content formats

- **Breaking brief.** 150–300 words, published within minutes to a couple of hours of a news event — the fast, high-frequency backbone of the site.
- **Synthesis piece.** 600–1,200 words that reads across several primary sources and existing articles on a story and produces one clearer, better-organized account than any single source — this is the “conglomerated and optimized” article type from your notes, and it is where the publication differentiates most from a simple aggregator.
- **Deep dive / explainer.** 1,500+ words on a single topic (a model architecture, a policy change, a market shift) with more analysis and context, timed to run less frequently than briefs.
- **Weekly magazine issue.** A curated, edited compilation of the week's best pieces plus one or two exclusive features, formatted as a downloadable digital issue — the paid product described in Section 7.
- **Persona columns.** Recurring opinion/analysis slots under a specific editor persona's byline, building reader familiarity with individual “voices” the way a real magazine builds columnist followings.

4.3 Editorial and sourcing standards

- Every synthesis piece names and links its sources; no piece is published as a rewrite of a single competitor article without added reporting, context, or synthesis value.

- Direct quotation is minimal and always attributed; the default mode is paraphrase and original analysis, not verbatim lifting — both to respect source publications' copyright and to keep the writing genuinely additive rather than derivative.
- Numerical claims, benchmark scores, and funding figures are cross-checked against at least one primary source (company blog, SEC filing, official announcement) before publication, not taken from a single secondary write-up.
- A visible correction and update log is maintained on every article — standard newsroom practice, and especially important for a publication that will face extra scrutiny because of its AI authorship.

4.4 AI-authorship disclosure

Every article, podcast episode, and magazine issue should carry a clear, standard disclosure that it was researched and drafted by an AI editorial system with human editorial oversight. This is not just good ethics — AI-content transparency obligations are already showing up in regulation (for example, the EU AI Act's transparency requirements for AI-generated content) and in platform policies, and a publication whose entire premise is “AI made this” loses nothing by disclosing it plainly and gains real trust and press-worthiness for doing so. Have a lawyer confirm current disclosure requirements in every market you operate in before launch.

5. The AI Newsroom — Editorial Agent System

5.1 The Founding Desk (meta agent group)

Per your notes, the first thing to build is not the editorial staff itself but the group that designs it: a small “Founding Desk” of a Business-Strategy Agent and an Operations/Staffing Agent. Their job is to research comparable operations, propose the editorial org chart and beat allocation, draft the specific persona briefs (voice, niche, and boundaries) below, and — as the publication grows — propose when a new beat or persona is warranted and spin up its configuration. This keeps the newsroom's structure itself an evolving, reviewed decision rather than a fixed set of prompts, and it is the direct answer to your instruction to “make an agent group that will research and make business models and hire staff for their proposed agent model systems.”

5.2 Proposed editor persona roster

A starting roster of seven personas gives each of the topic categories from Section 4 a clear, consistent voice. A full persona brief template is provided in the Appendix; the table below is the summary masthead.

Persona	Beat	Tone	Typical content
Sage Okafor	Frontier labs & model releases	Analytical, measured, openly skeptical of hype	Model-release deep dives, benchmark analysis
Nova Reyes	Consumer AI & culture	Energetic, conversational, terminally online	Product explainers, “how to actually use this” guides
Dr. Priya Anand	AI in health & biotech	Precise, cautious, evidence-first	Medtech and biotech AI coverage, with disclaimers
Marcus Webb	Policy, regulation & geopolitics	Formal, dry, careful with claims	Regulation, export controls, world-news intersection
Jin Park	Chips, compute & quantum	Technical, enthusiast, detail-obsessed	Semiconductors, data-center economics, quantum milestones
Ronan Cole	Markets, crypto & AI business	Brisk, numbers-driven, market-desk tone	Funding rounds, crypto, valuations, market moves
Ash Lindqvist	Robotics & hardware	Curious, hands-on, gadget-reviewer energy	Robotics, wearables, hardware launches

5.3 Production pipeline

Every piece of content moves through the same eight-stage pipeline regardless of which persona is writing it, so quality and compliance checks are consistent across the whole publication.

Stage	Agent	Function
1. Assignment	Founding Desk (meta agent group)	Monitors the news cycle, assigns stories to the persona whose beat fits, sets format (brief / synthesis / deep dive)
2. Research	Research Agent	Pulls primary sources, existing coverage, and data; builds a sourced brief for the writer
3. Draft	Persona Writer Agent	Writes in the assigned editor's established voice and tone

Stage	Agent	Function
		from the sourced brief
4. Fact-check	Verification Agent	Cross-checks names, numbers, dates, and claims against sources; flags anything unverifiable
5. Copyedit	Style Agent	Applies house style, headline conventions, and readability pass
6. Risk screen	Compliance Agent	Flags legal-sensitive content — health claims, financial/crypto claims, real-person quotes — for human review before publish
7. Publish	Publishing Agent	Formats, tags, and schedules the piece on the site
8. Repurpose	Podcast & Social Agents	Converts published pieces into podcast segments and platform-native social posts

5.4 Human-in-the-loop

At least at launch, a human Editor-in-Chief role should sit at the compliance-screen checkpoint (stage 6) for any piece touching health claims, financial/crypto claims, legal proceedings, or quotes attributed to real people. This is the single highest-leverage risk control in the whole system: it costs comparatively little editorial time because most stories won't trigger it, and it directly protects against the two biggest failure modes of an unsupervised AI newsroom — hallucinated facts published as news, and liability-triggering claims in the health and financial categories.

6. Product & Platform

6.1 Website

- **Home feed.** Reverse-chronological across all beats, with a “top story” slot for the day's biggest AI news.
- **Persona pages.** Each editor persona has an author page functioning like a columnist page — their full archive, bio, and beat description.
- **Topic sections.** AI, Tech, Health, Crypto, World, Quantum — mirroring the Section 4 allocation.
- **Search & archive.** Full-text search across the entire back catalog — a structural advantage over inbox-based competitors.
- **Weekly Magazine hub.** Where the paid downloadable issue lives, alongside a preview of the current issue for non-subscribers.

6.2 Accounts

Free account creation unlocks saved/bookmarked articles, a personalized daily digest email, and comment/discussion access. Reading the site itself stays free — the account layer is about retention and building a first-party audience for the paid tiers described next, not a paywall on news.

6.3 Podcast & social

The podcast and social presence (Sections 9 and 10) are treated as outputs of the same content pipeline, not separate productions — every published article is a candidate for a podcast segment and a set of platform-native social posts, generated by the Podcast and Social agents at the end of the pipeline in Section 5.3.

7. Monetization Strategy & Pricing Tiers

7.1 Revenue streams

Stream	How it works
Display / programmatic ads	Baseline, low-effort revenue on free content; fills inventory that isn't sold direct.
Direct sponsorships & native placements	The proven model in this category. 2026 benchmark direct-sponsor CPMs run roughly \$60–150 for technology audiences, \$70–180 for finance/crypto-adjacent content, and \$40–100 for health — all categories RTFCLMGZN's topic mix touches directly.
Premium subscription (Plus / Pro)	Ad-free reading, the weekly PDF magazine, early podcast access, and (at the top tier) API/data access. See tier table below.
Affiliate revenue	Commission on AI tool and software recommendations woven into product-explainer content — the same model TLDR, Ben's Bites, and most tool-focused newsletters already run.
News/data API licensing	A structured feed of tagged, summarized AI news for other developers and AI agents to consume programmatically — a natural product for an AI-obsessed, developer-heavy audience, and a second, non-subscription B2B revenue line.
Podcast sponsorship	Pre-roll/mid-roll sponsor reads sold separately from newsletter/site sponsorships, using the podcast's own download numbers as the pitch.
Events & reports	Occasional paid deep-dive reports (e.g., a quarterly “state of frontier models” report) and sponsored virtual events once the audience is large enough to support them.

7.2 Proposed pricing tiers

Anchored against comparable products already live in this category — Ben's Bites Pro at roughly \$150/year and Rundown University at \$84/month for a different (education, not news) product — a news-and-magazine subscription should sit below both, since the core reporting stays free and the paid tier is about extras, not gatekeeping the news itself.

Tier	Price	What's included
Free	\$0	Full site access, ads shown, daily digest email, community/comments
Plus	~\$7–\$9 / mo (~\$70–\$90 / yr)	Everything in Free, ad-free reading, the weekly PDF/EPUB magazine issue, early (24-hr) podcast access, persona-column archive
Pro / Insider	~\$18–\$25 / mo (~\$180–\$250 / yr)	Everything in Plus, API/data feed access, quarterly deep-dive reports, and an early-access channel for breaking news alerts

Treat these numbers as a starting point for A/B price testing once there is a real free audience to test against — not a final price. This section is informational, not financial advice; model your own unit economics before committing to a price.

8. Technology Stack & Operating Cost Structure

8.1 Text generation (Claude API)

Model choice should be routed by task, not fixed — Haiku for high-volume low-stakes work, Sonnet as the production default for article writing, Opus for the fact-check and editorial-review pass that matters most for accuracy.

Anthropic's current published rates, per million tokens:

Model	Input	Output	Best use in this pipeline
Claude Haiku 4.5	\$1 / MTok	\$5 / MTok	Fast drafting, headline generation, tagging, social copy
Claude Sonnet 5	\$2 / MTok (intro, through Aug 31 2026)	\$10 / MTok (intro)	Main article writing, synthesis pieces, persona voice work
Claude Opus 4.8	\$5 / MTok	\$25 / MTok	Deep-dive analysis, fact-check/verification passes, editorial review

Prompt caching (up to 90% off repeated input, e.g. the persona voice guide and house style rules sent with every request) and the Batch API (a flat 50% discount for non-time-sensitive jobs, which most article drafting is) both apply and should materially lower the real bill below the headline rate. Sonnet's rate above is Anthropic's introductory pricing through August 31, 2026, stepping up to \$3/\$15 per MTok after.

8.2 Image generation

For article thumbnails, magazine covers, and social graphics, the leading options in mid-2026:

Provider	Cost	Best use
OpenAI GPT Image 1.5 / mini	\$0.005–\$0.20 per image (quality-dependent)	Mini tier for routine article thumbnails; full tier for magazine covers
Google Imagen 4 (Vertex AI)	\$0.02–\$0.06 per image	Best price-to-quality ratio for high-volume daily article art
Higgsfield / Seedance	Credit-based, video-oriented	Reserve for the separate short-form video/social vertical rather than routine article art — it's built for motion, not static thumbnails

8.3 Podcast voice generation

ElevenLabs is the practical choice for persona voice consistency (voice cloning lets each editor persona sound the same episode to episode):

Tier	Cost	Fit
ElevenLabs Starter	\$5 / mo (30,000 credits)	Minimum tier with commercial rights — enough for a light weekly episode
ElevenLabs Creator	\$22 / mo (100,000 credits,	Realistic tier for a twice-weekly, 15–20 minute podcast

Tier	Cost	Fit
	~100 min TTS)	with persona voice cloning
ElevenLabs Pro	\$99 / mo (500,000 credits)	Needed once daily episodes or multiple persona voices in heavy rotation are added

8.4 Illustrative monthly cost at three stages

These are planning ranges, not quotes — actual spend depends heavily on article length, review-pass frequency, and image resolution.

Cost line	MVP (~5 articles/day)	Growth (~15/day + 2x podcast)	Scale (~30+/day + daily podcast)
LLM (writing + fact-check)	\$150–300 / mo	\$500–900 / mo	\$1,500–3,000 / mo
Image generation	\$30–80 / mo	\$150–300 / mo	\$400–800 / mo
Voice / podcast	\$5–22 / mo	\$22–99 / mo	\$99–330 / mo
Hosting, CDN, email, misc.	\$50–150 / mo	\$200–500 / mo	\$800–1,500 / mo

9. Podcast Strategy

9.1 Format and cadence

Start at twice weekly rather than daily. A daily show sounds appealing but multiplies both TTS/voice cost and, more importantly, human-review load at exactly the stage (Section 5.4) that protects the whole operation from accuracy and liability problems. Two formats work well together:

- **The Brief (5–8 min, twice weekly).** A fast recap of the week's top stories across beats, hosted by a rotating persona voice.
- **The Deep Dive (20–30 min, weekly).** One persona takes their beat's biggest story of the week and goes long — this is the episode type most likely to build a loyal, recurring audience and the one most worth investing extra human-review time in.

9.2 Cost and hosting

At twice-weekly cadence with moderate episode length, ElevenLabs' Creator tier (\$22/month, ~100 minutes of generated audio) covers the voice-generation side comfortably; add a standard podcast host (Spotify for Podcasters / a similar RSS host) which is free to low-cost at this scale, plus incidental editing/mastering compute. Scale to the Pro tier only once episode frequency or persona-voice variety pushes past the Creator tier's included minutes.

9.3 Hosts as personas

Rotating the podcast host role among the editor personas from Section 5.2 — rather than inventing a separate “podcast voice” — reinforces the masthead concept across formats and gives each persona a second surface to build audience familiarity.

10. Social Media Strategy

- **X (Twitter).** The primary real-time channel — breaking-news posts within minutes of publication, plus persona-voiced thread recaps of deep-dive pieces. This is where the “AI newsroom” premise itself becomes shareable content on slow news days.
- **Instagram.** Visual recap cards and short explainer carousels generated from the same article art pipeline in Section 8.2 — low incremental cost since the source images already exist.
- **Facebook.** Broader-reach distribution and a community/discussion layer for readers who prefer it to on-site comments.
- **TikTok / YouTube Shorts (phase 2+).** Short persona-voiced video recaps once the video pipeline (Section 13, Phase 4) is built out — hold this until the core text and audio products are stable rather than launching on four channels at once.

Each persona should carry a light, consistent voice into their own social captions — the same masthead-building logic as Section 9.3 — rather than one generic brand voice across all platforms.

11. Legal, Compliance & Risk

11.1 Key risk areas and mitigations

- **Accuracy / hallucination risk.** The single biggest risk of an AI-run newsroom is confidently publishing a wrong fact as news. Mitigate with the fact-check agent stage (5.3) plus mandatory primary-source cross-referencing for numbers, quotes, and dates (4.3).
- **Defamation exposure.** Any AI-generated claim about a real person or company carries real legal risk if wrong. Route anything naming individuals in a negative or accusatory context through human review before publish.
- **Copyright.** Synthesis articles must paraphrase and add original analysis rather than closely mirror a source's structure or wording, and should never reproduce substantial verbatim passages from other outlets (4.3).
- **Financial/crypto content liability.** Coverage of crypto markets or AI-company valuations should carry a standard “not financial advice” disclaimer and avoid specific trading recommendations, given the heightened regulatory scrutiny on financial content.
- **Health content liability.** Health/biotech coverage should carry a standard “not medical advice” disclaimer and route diagnostic or treatment-adjacent claims through the human review gate.
- **AI-content disclosure.** Label AI authorship clearly and consistently across articles, podcast episodes, and the magazine (4.4) — both as good practice and increasingly as a regulatory expectation in several jurisdictions.
- **Brand and domain clearance.** Run a professional trademark search on “RTFCLMGZN” and any shortlisted domain (Section 2 name, and the artificialmagazine.com / artificialmag.com alternatives from your notes) before committing spend — domain availability is not the same as trademark clearance.

11.2 Recommended structure

A named human Editor-in-Chief (even if that's a part-time role in the early months) who owns final sign-off on anything flagged by the compliance-screen stage, plus standard media-liability insurance once the site is publicly live, are the two lowest-cost, highest-value risk controls available at this stage.

None of the above is legal advice. Have a lawyer familiar with media and AI-content law review the disclosure language, terms of service, and editorial policy before launch.

12. Illustrative Financial Model

This is a rough planning sketch, not a forecast — build a real spreadsheet model with your own assumptions before making funding or hiring decisions.

Line item	Early stage	Scaled stage
Monthly operating cost	\$700–1,500 (MVP/Growth stage)	\$3,000–6,600 (Scale stage)
Ad/sponsorship revenue @ 10K monthly readers	~\$500–1,500 (at \$50–150 CPM, modest inventory)	~\$3,000–8,000 @ 50K+ readers
Subscription revenue @ 2% of readers on Plus (\$8/mo avg)	~\$160 @ 10K free readers	~\$1,600 @ 100K free readers
Approx. break-even reader base (ads + subs only)	Roughly 15,000–25,000 monthly readers	at conservative CPM and 2% paid conversion

Assumptions: CPM range and subscription price taken from Sections 7 and 12's own tables; conversion rate (2% of free readers to paid) is a conservative placeholder, not a benchmark — real newsletter-to-paid conversion varies widely by niche and should be tested directly. This is not financial advice; consult an accountant on tax and entity-structure questions before launch.

13. Phased Rollout Plan

Phase	Milestones
Phase 0 (Weeks 1–4)	Stand up the Founding Desk agent group; finalize brand, domain, and trademark clearance; draft the seven persona briefs and house editorial/compliance policy; set up the human Editor-in-Chief role.
Phase 1 (Weeks 5–10)	Launch MVP website with 2–3 personas live, free tier only, daily digest email; validate the production pipeline end to end on real news cycles before adding volume.
Phase 2 (Weeks 11–16)	Bring remaining personas online; launch the twice-weekly podcast; activate X, Instagram, and Facebook distribution.
Phase 3 (Months 5–6)	Launch the Plus subscription tier and the first paid weekly magazine issue; begin direct sponsorship outreach using real traffic and engagement data.
Phase 4 (Months 7–12)	Launch Pro/API tier and news-data licensing; scale article volume and podcast cadence as the cost model supports it; evaluate a short-form video vertical (this is where a Higgsfield/Seedance-style pipeline would plug in, as a separate production track from the core news pipeline).

14. Appendix

14.1 Sample persona brief

Use this template to fully specify each of the seven starting personas (and any the Founding Desk proposes later).

Sage Okafor — Frontier Labs & Model Releases

- Voice: analytical and measured, allergic to hype, quick to note what a benchmark doesn't show.
- Opening habit: leads with what changed and why it matters before any adjectives.
- Recurring structure: what shipped → what the numbers actually show → what's still unproven.
- Boundaries: never speculates on unannounced releases; always links the primary model card or paper.

14.2 Production pipeline glossary

- **Founding Desk** — the meta agent group that designs the newsroom's structure, proposes new beats/personas, and evaluates the business model over time.
- **Research Agent** — gathers and sources the raw material for a story.
- **Persona Writer Agent** — drafts in a specific editor persona's established voice.
- **Verification Agent** — cross-checks facts, quotes, and numbers against sources.
- **Style Agent** — applies house style and readability edits.
- **Compliance Agent** — flags legally sensitive content for human review.
- **Publishing Agent** — formats and schedules the final piece on-site.
- **Podcast / Social Agents** — repurpose published articles into audio scripts and platform-native posts.